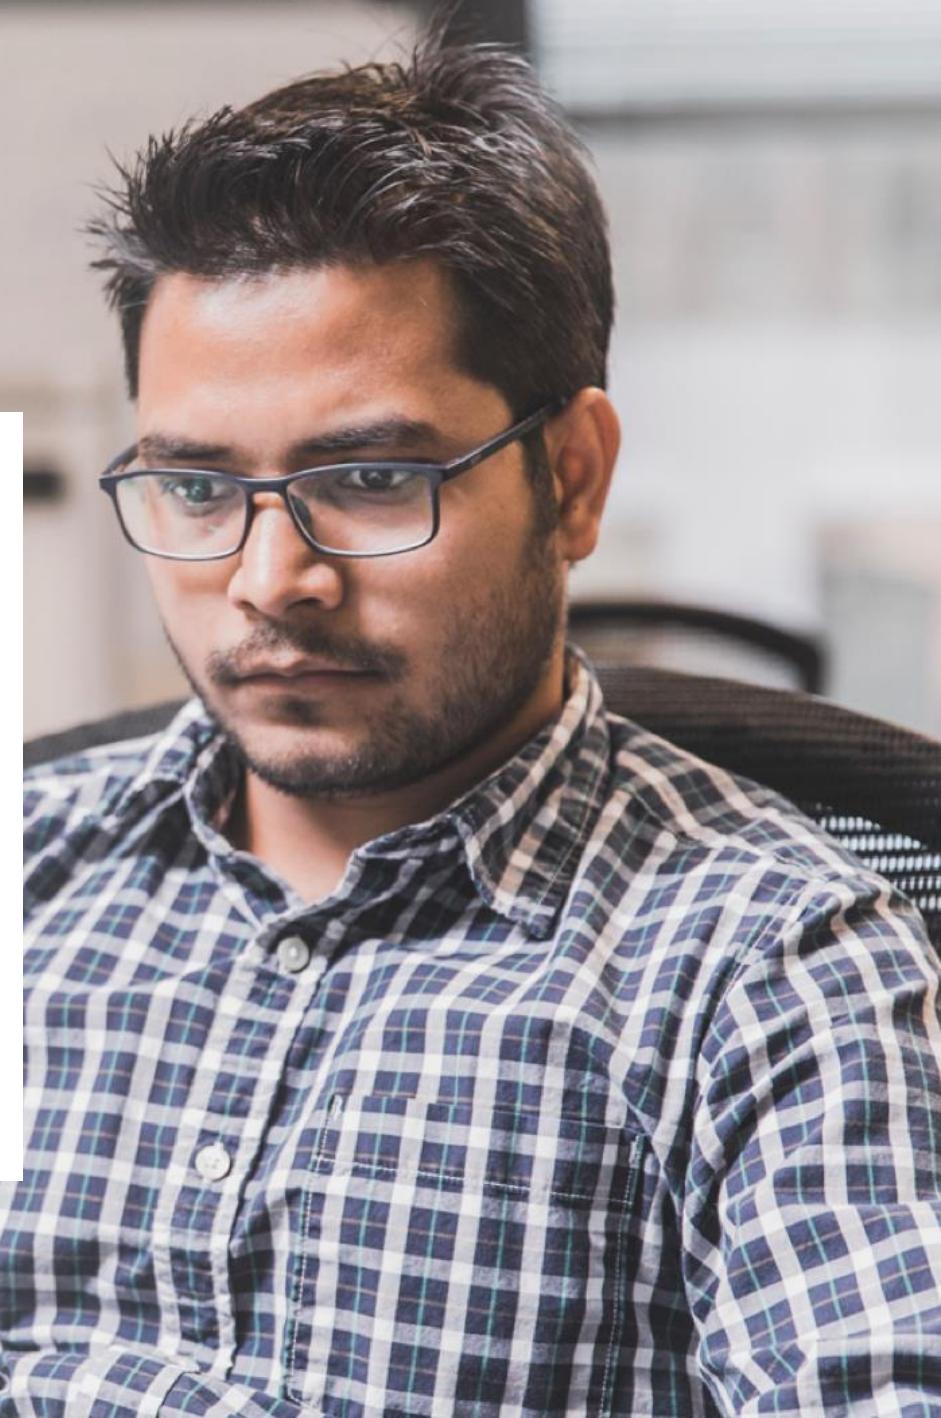


Azure Study Group

AZ-300 - Microsoft Azure Architect Technologies

Jeff Wagner
Partner Technology Strategist



Develop for the Cloud (20-25%)

Agenda

1

Series
Agenda

2

Speaker
Introduction

3

Feedback
Loop

4

Objective
Review

5

Open Mic

Series Agenda

- | | |
|---|--|
| 1 | Deploy and Configure Infrastructure (25-30%) |
| 2 | Implement Workloads and Security (20-25%) |
| 3 | Create and Deploy Apps (5-10%) |
| 4 | Implement Authentication and Secure Data (5-10%) |
| 5 | Develop for the Cloud (20-25%) |
- <https://aka.ms/azurecsg>

Series Agenda

- | | |
|---|--|
| 1 | Deploy and Configure Infrastructure (25-30%) |
| 2 | Implement Workloads and Security (20-25%) |
| 3 | Create and Deploy Apps (5-10%) |
| 4 | Implement Authentication and Secure Data (5-10%) |
| 5 | Develop for the Cloud (20-25%) |
- <https://aka.ms/azurecsg>

Speaker Introduction - Jeff Wagner

- Partner Technology Strategist based in Atlanta
- 21+ years with Microsoft, more in the industry
- Been working with Microsoft Azure when we weren't sure if it was called Windows Azure or Windows *Azure*
- Constant learner - *Ancora Imparo*



Feedback Loop

Objectives

Configure a message-based integration architecture

May include but not limited to: Configure an app or service to send emails, Event Grid, and the Azure Relay Service; create and configure Notification Hub, Event Hub, and Service Bus; configure queries across multiple products

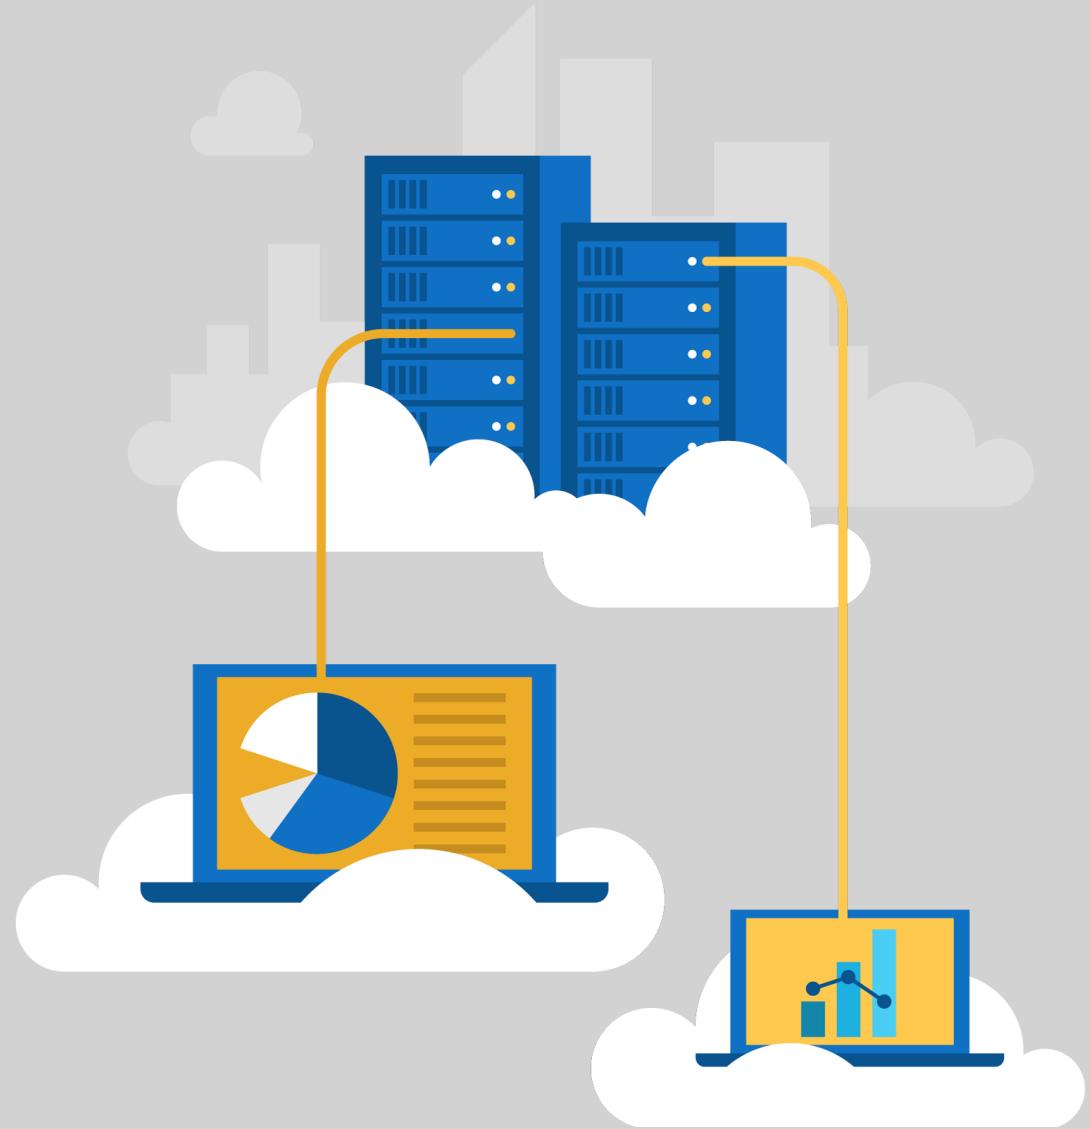
Develop for autoscaling

May include but not limited to: Implement autoscaling rules and patterns (schedule, operational/system metrics, code that addresses singleton application instances); implement code that addresses transient state

Configuring a Message-Based Integration Architecture



Configure an app or service to send emails



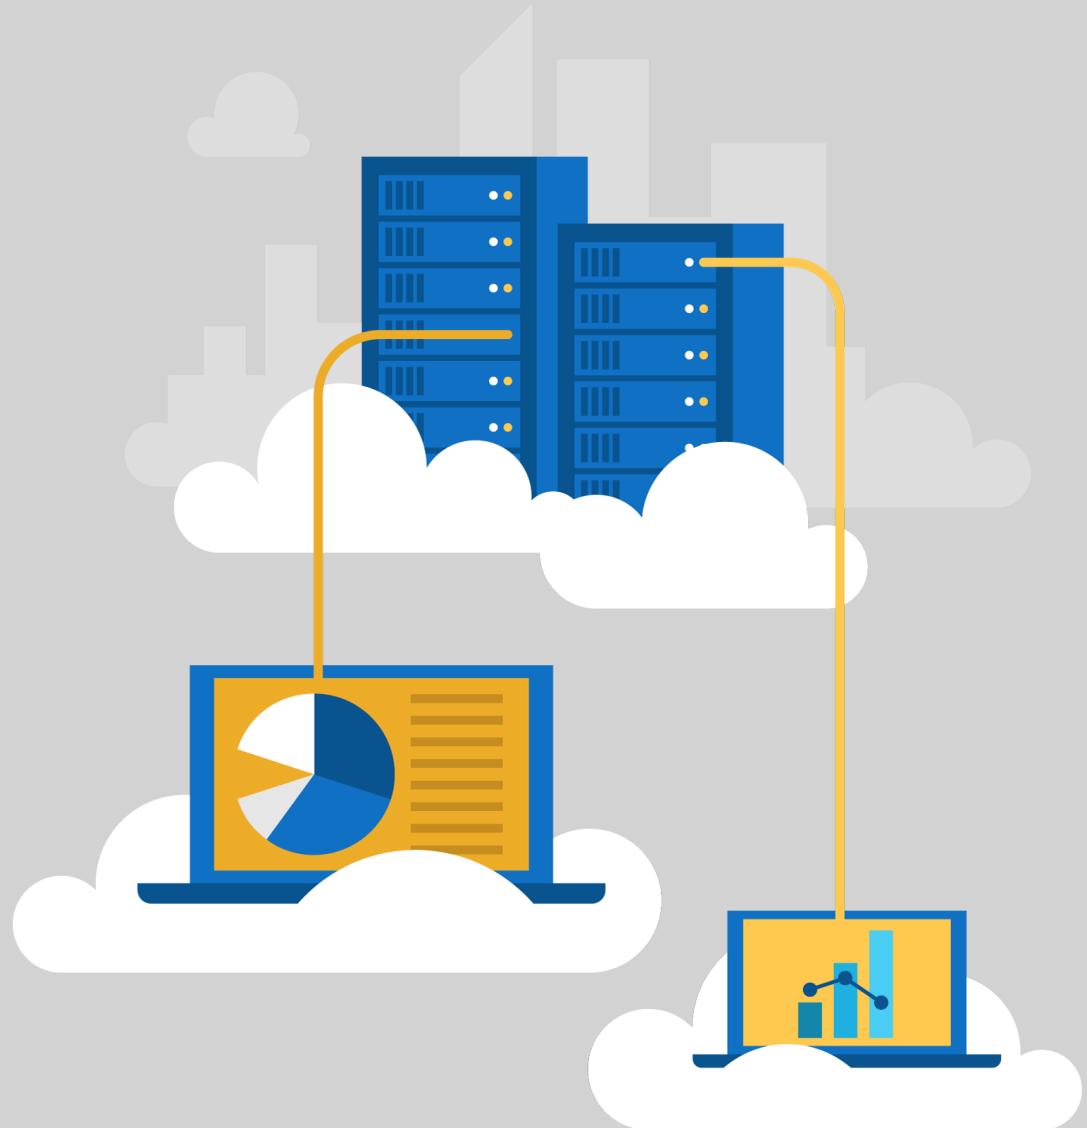
SendGrid

- A third party, cloud-based messaging service, which offers:
 - **deep level of integration with the Azure platform and the Azure portal**
 - **transactional email delivery**
 - **scalability based on email volume**
 - **real-time analytics**
 - **a flexible API to enable custom integration scenarios.**

Sending emails through SendGrid by using C#

- Based on an open source library:
 - Hosted on GitHub at <https://github.com/sendgrid/sendgrid-csharp>
 - Available as a NuGet package
- Implementation steps:
 - Create a new instance of the **SendGridClient** class
 - Pass a string parameter containing API Key
 - Use class methods to perform common tasks:
 - **SendGridClient.SendEmailAsync** to send a simple email
 - **AddAttachment** to attach a file to an email message

Configure an event publish and subscribe model



Event-driven architecture

- Consists of event producers and event consumers:
 - **Facilitates processing of events in near-real time**
 - **Decouples producers from consumers**
 - **Decouples consumers from each other**
 - **Allows all consumers to see all events**
- Offers several Azure-based implementations:
 - **Simple event processing:** Each event triggers immediately an action in a consumer:
 - e.g. Azure Functions with an Azure Service Bus trigger
 - **Complex event processing:** Consumers process a series of events, looking for patterns:
 - e.g. Azure Stream Analytics or Apache Storm
 - **Event stream processing:** A stream processor processes or transforms the stream:
 - e.g. Azure IoT Hub or Apache Kafka

Azure Event Grid

- Simplifies building applications with event-based architectures:
 - Reduces implementation to:
 - Selecting an Azure resource to subscribe to
 - Providing the event handler or a webhook endpoint to send the event to
 - Offers built-in support for events originating from a range of Azure services, including:
 - Azure subscription management operations, custom topics, Azure Event Hubs, IoT Hub, Azure Media Services, Resource group management operations, Service Bus, Azure Blob storage, General Purpose v2 storage
 - Integrates with a range of Azure services that provide event handling:
 - Azure Automation, Azure Functions, Event Hubs, Hybrid Connections, Azure Logic Apps, Microsoft Flow, Azure Queue storage

Subscribing to Blob storage events using Azure CLI

- Implementation steps:
 - **Create an Azure resource group:**
 - `az group create --name DemoGroup --location eastus`
 - **Create an Azure Storage account (Blob storage or General Purpose):**
 - `az storage account create --name demostor --location eastus --resource-group DemoGroup --sku \ Standard_LRS --kind BlobStorage --access-tier Hot`
 - **Identify the resource ID of the storage account:**
 - `storageid=$(az storage account show --name demostor --resource-group DemoGroup \ --query id --output tsv)`
 - **Create a subscription:**
 - `az eventgrid event-subscription create --resource-id \$storageid --name \ contosostoragesub --endpoint https://contoso.com/api/updates`

Configure the Azure Relay service



Azure Relay service

- Facilitates hybrid connectivity:
 - **Securely connects public cloud to services within a corporate enterprise network**
 - **Eliminates the need to open inbound firewall ports or change network infrastructure**
 - **Uses an outbound port to establish bidirectional socket dedicated to an internal service**
- Supports a variety of connectivity scenarios:
 - **Traditional one-way, request/response and peer-to-peer traffic**
 - **Event distribution to enable publish/subscribe scenarios**
 - **Bidirectional socket communication**
- Offers the Hybrid Connections capability, which additionally supports:
 - **WebSocket protocol**
 - **HTTP and HTTPs requests and responses**

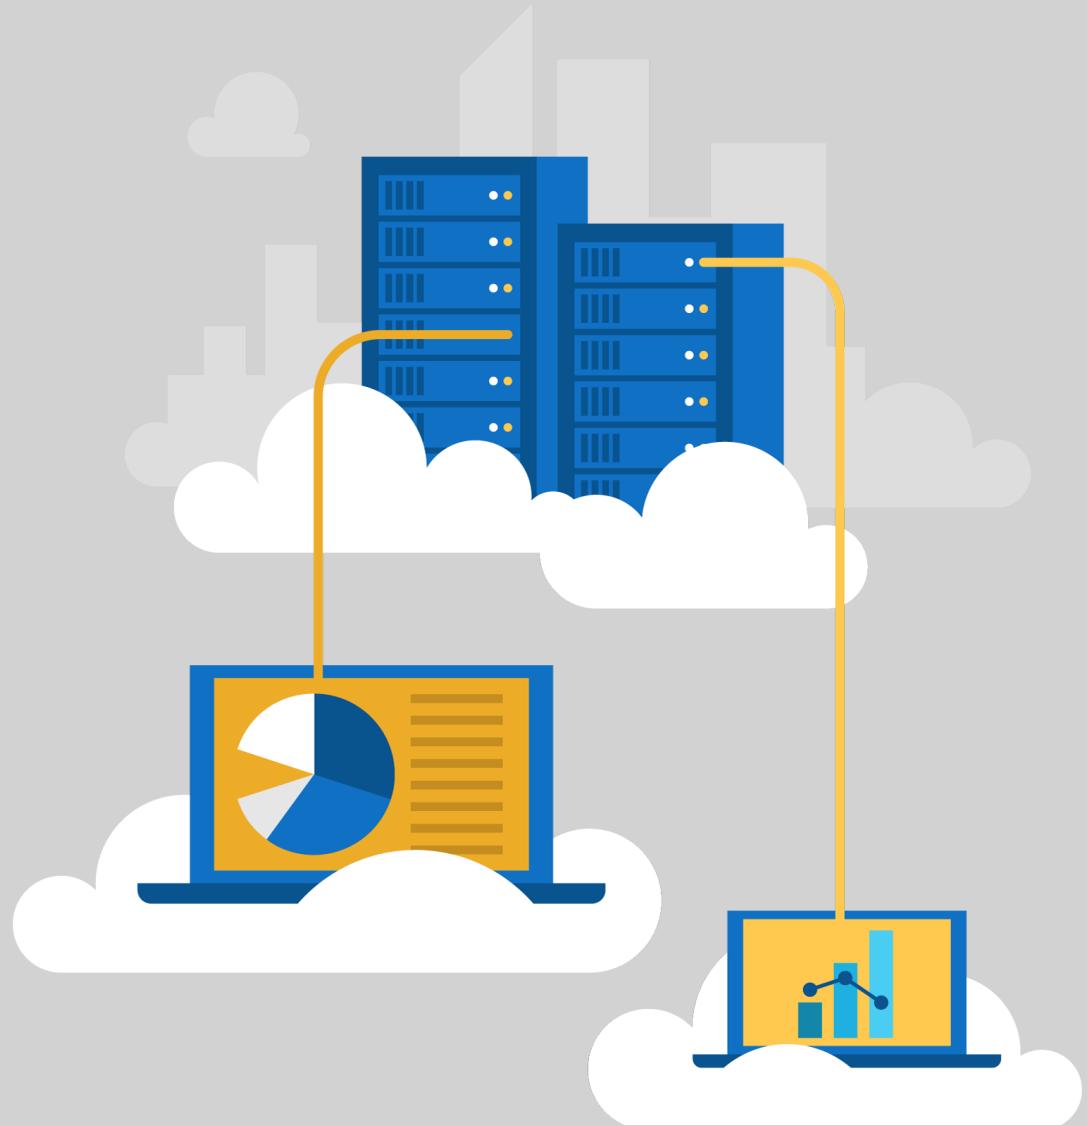
Azure Relay service architecture

- An Azure load balancer: routes requests to any of the gateway nodes.
- Gateway nodes: create new relays, forward connection requests to specific nodes that own a requested relays, and send rendezvous requests to listening clients.
- Listening clients: create temporary channels to gateway nodes in response to rendezvous requests and exchange messages via the gateway node.
- Communication flow:
 - A client (Client A) creates a listening request that is routed to a gateway.
 - The gateway creates a new relay.
 - A different client (Client B) creates a connection request.
 - The connection request from Client B is handled first by looking up the associated relay.
 - Once the correct relay has been identified, the request is forwarded to that specific relay.
 - A rendezvous request is sent to Client A.
 - Client A creates a temporary channel to Client B via the gateway that Client B used in its connection request.
 - Client B can receive messages from Client A directly from its original gateway.
 - Client A can receive messages from Client B via the original gateway and the established temporary channel.

Using Azure Relay in Node.js

- Facilitated by:
 - **Node.js ws package WebSocket protocol client library**
 - **Node.js hyco-ws package for Hybrid Connections in Azure Relay, which extends ws:**
 - Adds a new server class exported via `require('hyco-ws').RelayedServer` and a few helper methods
 - Provides `hycows.RelayedServer` class as an alternative to the `ws.Server` class. The `RelayedServer` constructor has two required arguments to establish a connection over the WebSocket protocol using Azure Relay:
 - `server` - The fully qualified URI for a Hybrid Connection name on which to listen.
 - `token` - Either a previously issued token string or a callback function that can be called to obtain such a token string.
 - **Any libraries that support the WebSocket protocol in the language of your choice.**

Create and configure a notification hub



Azure Notification Hubs

- A scaled-out push engine that:
 - **Allows you to send notifications:**
 - to practically any platform (iOS, Android, Windows, Kindle, BAIDU, etc.)
 - from practically any back end (cloud or on-premises).
 - **Eliminates the complexity associated with implementing push notifications:**
 - Offers multi-platform, scaled-out push notification infrastructure
 - Devices are responsible only for registering their PNS handles with a hub.
 - **Facilitates a number of common notifications scenarios:**
 - Sending breaking news notifications to millions with low latency
 - Sending location-based coupons to interested user segments
 - Sending event-related notifications for media/sports/finance/gaming applications
 - Pushing promotional content to applications to engage and market to customers
 - Notifying users of enterprise events, like new messages and work items
 - Sending codes for multi-factor authentication

Configuring Notification Hubs in iOS

- Ensure that an Xcode application satisfies the following prerequisites:
 - **The Push Notifications capability must be enabled.**
 - **The files distributed in the Azure messaging framework must be included in the project.**
- Implement the following steps:
 - **add constants to HubInfo.h that will contain connection details for your notification hub**
 - **import the WindowsAzureMessaging/WindowsAzureMessaging and UserNotifications/UserNotifications directives into the project**
 - **add code to connect to the notification hub using the information stored in HubInfo.h**

Configuring Notification Hubs in Xamarin.Android

- Ensure that a Xamarin project satisfies the following prerequisites:
 - **The `Xamarin.GooglePlayServices.Base`, `Xamarin.Firebase.Messaging`, and `Xamarin.Azure.NotificationHubs.Android` NuGet packages must be installed.**
 - **The `google-services.json` file must be downloaded from the Google Firebase Console and then copied to the root of your project folder.**
 - **`com.google.firebaseio.iid.FirebaseInstanceIdReceiver` receiver must be registered to enable PNS registration and message receipt.**
- Implement the following steps:
 - **Add a C# class that will contain connection details for your notification hub**
 - **Create a C# class to manage PNS registration**
 - **Create a separate C# class to handle the receipt of a new message and display that message in the application's UI**

Sending messages from an application back end to Notification Hubs using C#

- Available via a number of client libraries:
 - In .NET applications, the **Microsoft.Azure.NotificationHubs** NuGet package includes the **NotificationHubClient** class that can be used to send messages
- Simplified by Azure Notification Hubs templates:
 - Allow you to specify how a device should receive notifications
 - Facilitate sending notifications to multiple platforms at the same time without having to specify a native payload

Create and configure an event hub



Azure Event Hubs

- A big data streaming platform and event ingestion service:
 - **Capable of receiving and processing millions of events per second**
 - **Facilitating processing and storing events, data, or telemetry produced by distributed software and devices**
 - **Providing low latency and seamless integration with data and analytics services inside and outside of Azure**
 - **Serving as an event ingestor in a pipeline between event producers and consumers**
 - **Consisting of the following main components:**
 - Event producers
 - Partitions
 - Consumer groups
 - Throughput units
 - Event receivers

Azure IoT Hub

- A managed service operating as a central message hub:
 - **Provides bidirectional communication between IoT application and IoT devices**
 - **Scales to millions of IoT devices of practically any type**
 - **Supports multiple messaging patterns, including:**
 - device-to-cloud telemetry
 - file upload from devices
 - request-reply methods
- Differs from Azure Event hubs in a number of aspects:
 - **IoT Hub:**
 - developed specifically to address the unique requirements of connecting IoT devices
 - includes features that enrich the relationship between IoT devices and back-end systems
 - **Event Hubs:**
 - designed for big data streaming
 - offer a partitioned consumer model to scale out stream and integrate with big data and analytics services

Create and configure a service bus



Azure Service Bus

- A messaging service that offers three communication mechanisms:
 - **Queues**, which allow one-directional communication. Each queue acts as an intermediary that stores sent messages until they are received. Each message is received by a single recipient.
 - **Topics**, which provide one-directional communication using subscriptions. A single topic can have multiple subscriptions. Like a queue, a topic acts as a broker, but each subscription can use a filter to receive only messages that match specific criteria.
 - **Relays**, which provide bidirectional communication. Unlike a queue or a topic, a relay doesn't store in-flight messages, it just passes them on to the destination application.
- Namespace is the scoping container for a Service Bus implementation:
 - **Multiple queues, topics, and relays can reside within a single namespace.**
 - **Namespaces often serve as application containers.**
 - **Each namespace name must be globally unique.**

Service Bus queues

- Support a brokered, asynchronous messaging communication model:
 - **A message producer (sender) pushes a message to a queue and continues its processing.**
 - **A message consumer (receiver) pulls the message from the queue and processes it.**
- Can be used for a wide variety of scenarios:
 - **Communication between web and worker roles in a multi-tier Azure application.**
 - **Communication between on-premises apps and Azure-hosted apps in a hybrid solution.**
 - **Communication between the components of a distributed application running on-premises in different organizations or departments of an organization.**

Manipulating a Service Bus queue using Ruby

- Implementation steps:
 - **Download the Azure Ruby package:**
 - `gem install azure`
 - **Add the `azure` directive at the top of Ruby file where you want to reference the Azure Ruby libraries:**
 - `require "azure"`
 - **Create a connection to Service Bus using the client object**
 - **Send messages to a Service Bus queue by calling the `send_queue_message()` method on the `Azure::ServiceBusService` object:**
 - **Messages are `Azure::ServiceBus::BrokeredMessage` objects and have:**
 - a set of standard properties (such as `label` and `time_to_live`)
 - a dictionary that is used to hold custom
 - application-specific properties
 - a body of arbitrary application data.

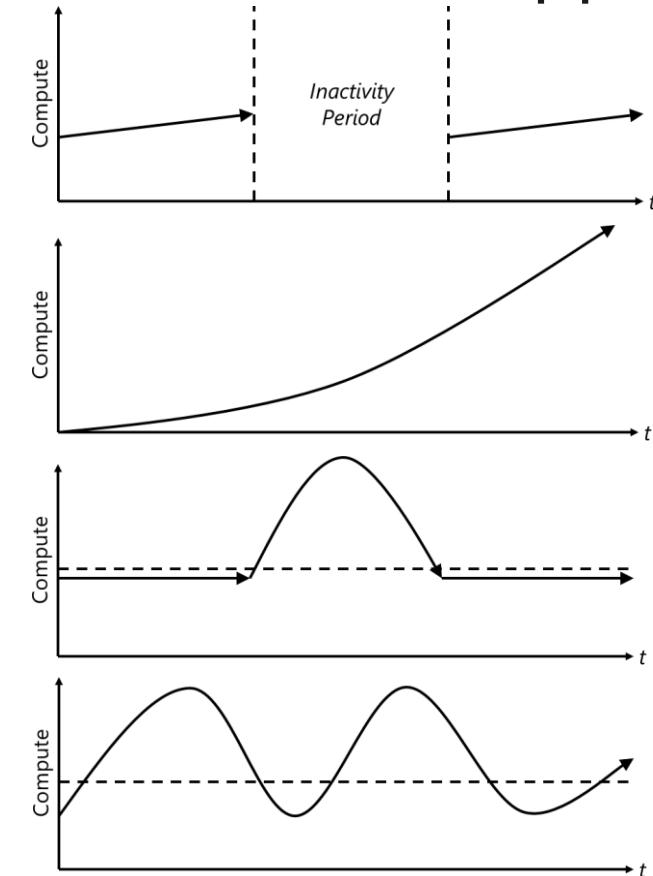
Developing for Autoscaling

Subtitle or speaker name



Computing patterns

- The most common computing patterns for cloud-based web apps:
 - **On and off:**
 - Common for batch processing
 - Waste of overprovisioned capacity
 - **Growing fast:**
 - Represents a successful service
 - Presents growth challenges
 - **Unpredictable bursting:**
 - Unexpected peaks of demand
 - Waste of overprovisioned capacity
 - **Predictable bursting:**
 - Requires resource provisioning and deprovisioning
 - Introduced increased management complexity



Scale and auto-scale

- Cloud offers elastic scaling by using its practically unlimited resources
- Auto-scaling is the ability to monitor workload usage and automatically scale according to changes in usage levels
- Many Azure services support manual and automatic scaling:
 - **infrastructure services such as Azure Virtual Machine Scale Sets**
 - **application services such as Azure App Service**
 - **database services such as Azure Cosmos DB.**

Auto-scale metrics

- App Service support autoscale:
 - Requires service plans using Basic, Standard, or Premium pricing tiers
 - Increases or decreases the instance count within the same service plan
 - Supports a number of auto-scale metrics:
 - CPU: CpuPercentage
 - Memory: MemoryPercentage
 - Data in: BytesReceived
 - Data out: BytesSent
 - HTTP queue: HttpQueueLength
 - Disk queue: DiskQueueLength
 - Allow you to trigger alerts when the value of a metric crosses a custom threshold:
 - Send email notifications to the service administrator and co-administrators
 - Send email to additional email addresses that you specify
 - Call a webhook
 - Start the execution of an Azure runbook

Implement code that
addresses singleton
application instances



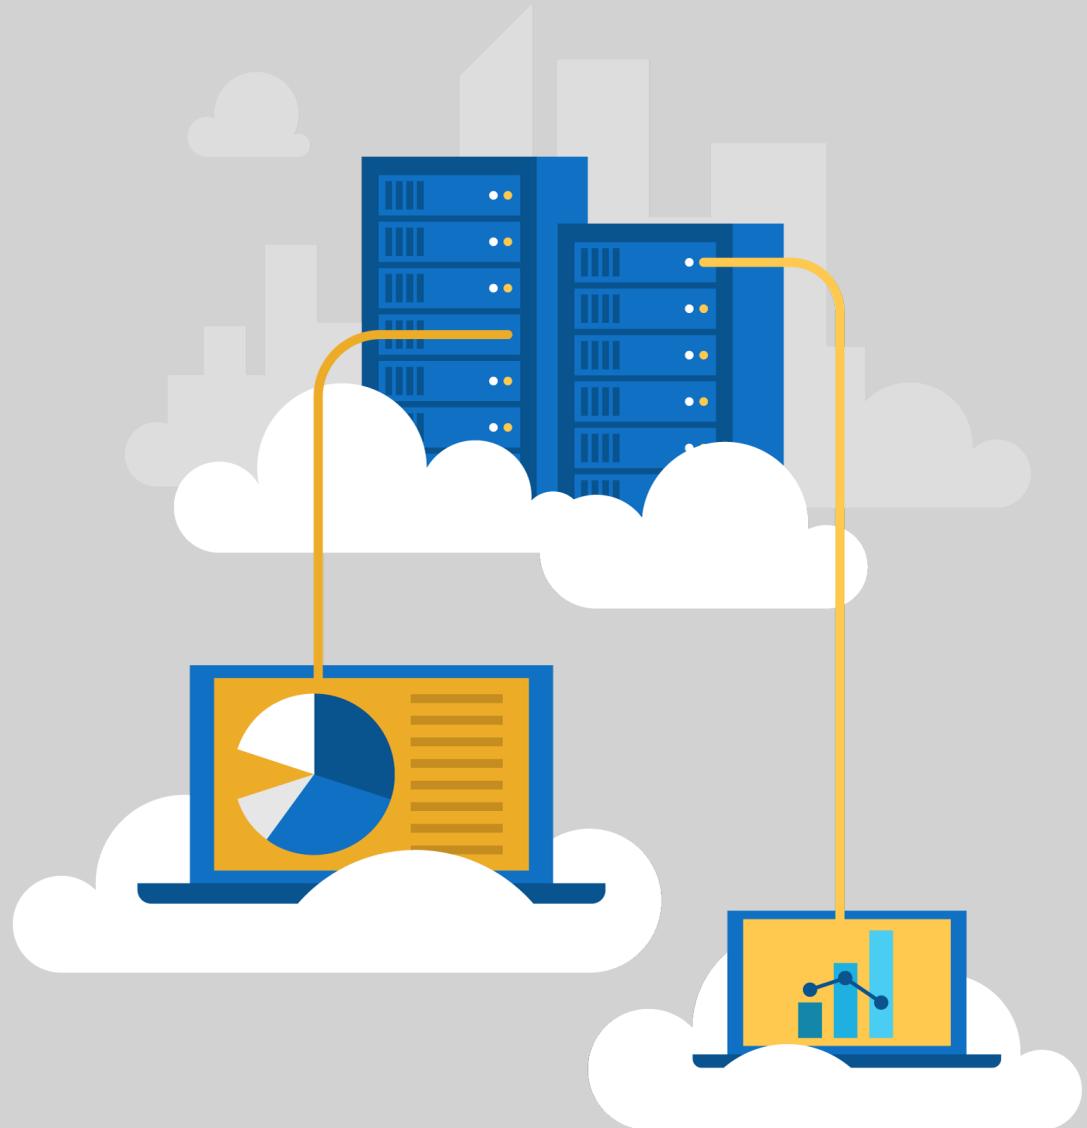
Querying resources using Azure CLI

- Azure CLI is optimized for managing Azure resources:
 - **Supports the --query argument to execute a JMESPath query on command results:**
 - A single result in the format of a JSON document
 - Multiple results in the format of a JSON array (you can flatten them by using the [] JMESPath operator)
 - **For example, to list Azure VMs in the current subscription, run:**
 - `az vm list`
 - **To limit results to the image name and offer, run:**
 - `az vm list --query '[].{name:name, image:storageProfile.imageReference.offer}'`
 - **To list Azure VMs with the Windows Server OS installed, run:**
 - `az vm list --query "[?starts_with(storageProfile.imageReference.offer, 'WindowsServer')]"`
 - **To list name and vmID or Azure VMs with Ubuntu OS installed, run:**
 - `az vm list --query "[?starts_with(storageProfile.imageReference.offer, 'Ubuntu')].{name:name, id:vmId}"`

Querying resources using the fluent Azure SDK

- Implementation steps:
 - **Create an authenticated client by invoking the Azure.Authenticate static method:**
 - The Authenticate method requires an authorization file necessary to access a target subscription:
 - `Azure azure = Azure.Authenticate("azure.auth").WithDefaultSubscription();`
 - You can generate the file by using the Azure CLI: `az ad sp create-for-rbac --sdk-auth > azure.auth`
 - **Query resources by using properties of the IAzure interface of the Azure variable:**
 - To list Azure VMs synchronously, use the VirtualMachines property:
 - `azure.VirtualMachines`
 - To list Azure VMs synchronously, use the VirtualMachines property:
 - `var vms = await azure.VirtualMachines.ListAsync();
foreach(var vm in vms)
{
 Console.WriteLine(vm.Name);
}`
 - To determine an IP address of an Azure VM, use the IVirtualMachine and related interfaces:
 - `INetworkInterface targetnic = targetvm.GetPrimaryNetworkInterface();`
 - `INicIPConfiguration targetipconfig = targetnic.PrimaryIPConfiguration;`
 - `IPublicIPAddress targetipaddress = targetipconfig.GetPublicIPAddress();`

**Implement code that
addresses a transient
state**

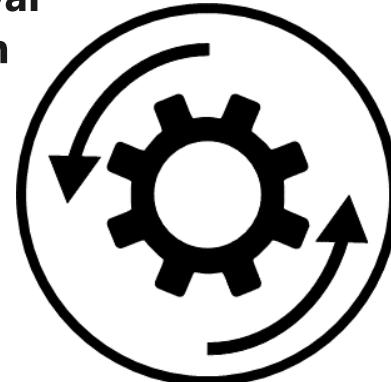


Transient errors

- Transient faults are typically self-correcting and include:
 - **momentary loss of network connectivity to components and services**
 - **temporary unavailability of a service**
 - **timeouts that occur when a service is busy**
- Applications that use cloud services should be able to handle gracefully transient faults:
 - **If the action that triggered a fault is repeated after a suitable delay, it is likely to succeed**

Handling transient errors

- Applications can handle a failure by using the following strategies:
 - **Cancel:** if the fault indicates that the failure is not transient
 - **Retry:** If the specific fault reported is unusual or rare
 - **Retry after delay:** If the fault is caused by a common connectivity or busy failure:
 - the period between retries should be chosen to result in an even distribution of requests from multiple instances of the application:
 - 1.The application invokes an operation on a hosted service. The request fails, and the service host responds with HTTP response code 500 (internal server error).
 - 2.The application waits for a short interval and tries again. The request fails with HTTP response code 500.
 - 3.The application waits for a longer interval and tries again. The request succeeds with HTTP response code 200 (OK).



Handling transient errors in code

- The `OperationWithBasicRetryAsync` method invokes an external service asynchronously through the `TransientOperationAsync` method.
- The details of the `TransientOperationAsync` method are service-specific
- The method is contained in a try/catch block wrapped in a for loop:
 - If the `TransientOperationAsync` method succeeds, the for loop exits.
 - If the `TransientOperationAsync` method fails, the catch block examines the reason for the failure. If it's believed to be a transient error, the code waits for a short delay before retrying the operation.
 - The for loop also tracks the number of times that the operation has been attempted.
 - If the code fails three times, the exception is assumed to be more long lasting.
 - If the exception is long lasting, the catch handler throws an exception.

Detecting if an error is transient in code

- The `IsTransient` method checks for a specific set of exceptions that are relevant to the environment the code is run in:
 - **Specifics are context-dependent**

```
private bool IsTransient(Exception ex)
{
    if (ex is OperationTransientException)
        return true;

    var webException = ex as WebException;
    if (webException != null)
    {
        return new[] {
            WebExceptionStatus.ConnectionClosed,
            WebExceptionStatus.Timeout,
            WebExceptionStatus.RequestCanceled
        }.Contains(webException.Status);
    }

    return false;
}
```



Questions?

Homework Assignment

GET EXAM READY!

Note the changes coming 2/22/19



Open Mic

